

Table of Contents

[aidocs](#)

[How It Works](#)

[Installation](#)

[Install from PyPI](#)

[Or install from GitHub](#)

[Or use pipx](#)

[Updating](#)

[1. Update the CLI](#)

[2. Reinstall commands in your project \(adds new slash commands\)](#)

[Quick Start](#)

[Install the CLI](#)

[Add to your project](#)

[Usage Flow](#)

[Simple: Generate docs for one page](#)

[Flow: Document a feature \(user-focused by default\)](#)

[Batch: Document entire project](#)

[Maintain: Update after code changes](#)

[RAG: Setup semantic search for your docs](#)

[CLI Commands](#)

[Slash Commands](#)

[How to Import Payments](#)

[Before You Start](#)

[Steps](#)

[What Happens Next](#)

[Import Payments Flow](#)

[Architecture](#)

[Entry Points](#)

Execution Flow

1. Generate migration
2. Run migration
3. Chunk your docs
4. Generate embeddings and sync

Knowledge Base

Intelligent Workflow

Option A: Document a specific flow

Option B: Document entire project

Resume if interrupted

After code changes

Configuration

Output

Requirements

Development

License

aidocs

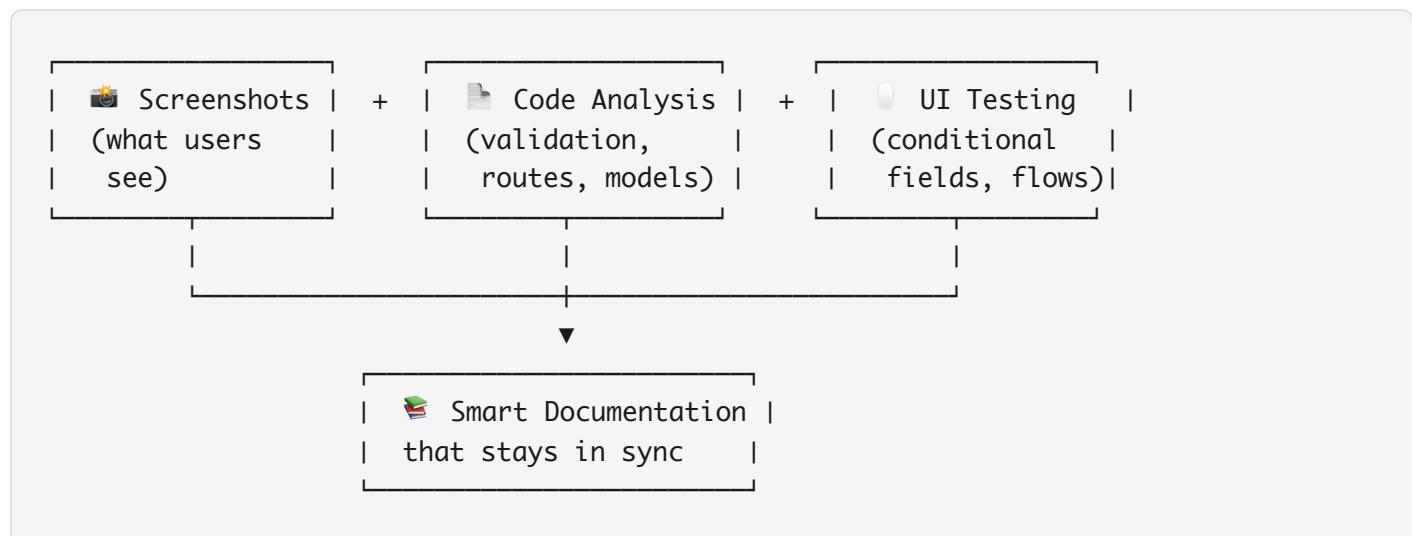
AI-powered documentation generator for web applications.

How It Works

aidocs generates comprehensive documentation by combining **three sources of truth**:

1. **Vision Analysis** - Playwright captures screenshots, Claude analyzes what users actually see
2. **Codebase Analysis** - Scans your frontend components, backend routes, validation rules, and models
3. **Interactive Exploration** - Clicks buttons, fills forms, discovers conditional UI and validation messages

This produces documentation that's accurate to both the code AND the actual user experience.



Installation

```
# Install from PyPI
uv tool install aidocs

# Or install from GitHub
uv tool install aidocs --from git+https://github.com/binarcode/aidocs-cli.git

# Or use pipx
pipx install aidocs
```

Updating

When a new version is released, update the CLI and reinstall commands in your project:

```
# 1. Update the CLI  
aidocs update  
  
# 2. Reinstall commands in your project (adds new slash commands)  
cd your-project  
aidocs init . --force
```

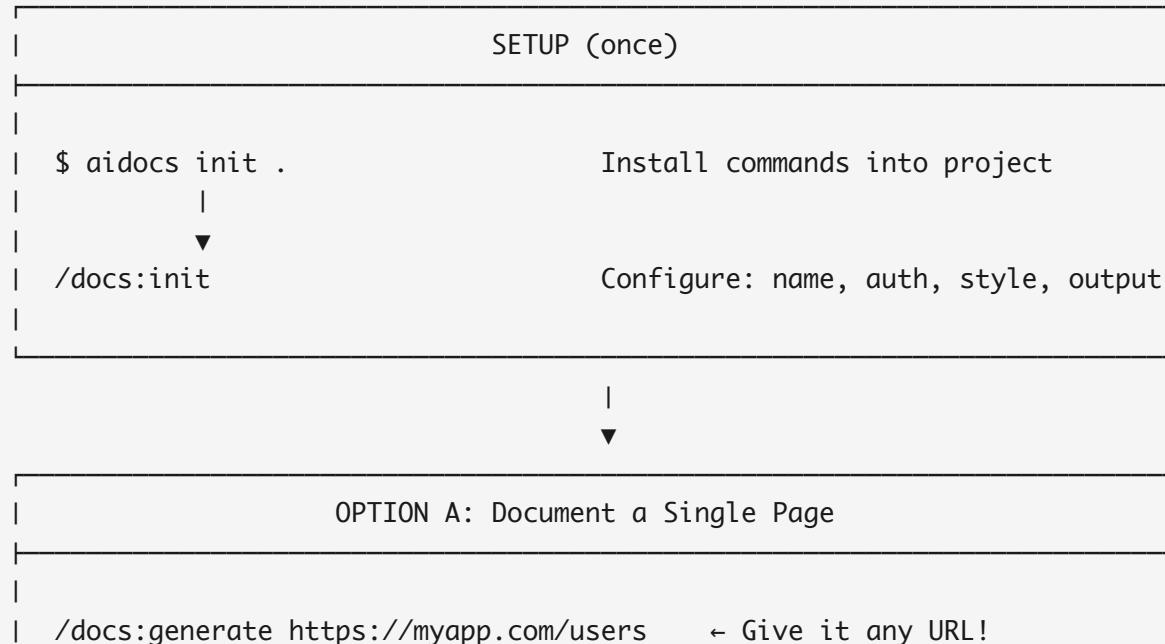
The `--force` flag overwrites existing command files, adding any new commands from the latest version.

Tip: Run `aidocs update --github` to get the latest unreleased features from GitHub.

Quick Start

```
# Install the CLI  
uv tool install aidocs  
  
# Add to your project  
aidocs init .
```

Usage Flow



```
|           |
|   └── Takes screenshots with Playwright
|   └── Analyzes codebase for that route
|   └── Documents UI elements and interactions
|       └── Creates docs/users/index.md
```



OPTION B: Document a Code Flow

```
/docs:flow "sync users from discord"      ← Describe the flow in words!
|
|   └── Searches codebase for relevant files
|   └── Traces execution path and builds call graph
|   └── Generates mermaid sequence diagram
|   └── Captures UI screenshots (if Playwright + route detected)
|       └── Creates docs/flows/sync-users-from-discord.md
```



OPTION C: Document Entire Project

```
/docs:discover          Scan codebase, find all modules
|
|   ↓
|
| /docs:plan             Create ordered documentation plan
|   |
|   ↓
| /docs:execute           Run through plan, generate all docs
|   |
|   → Resume with --continue if interrupted
```



KEEP DOCS IN SYNC

```
# After implementing a feature:
/docs:update --base main      Detect changes, update affected docs
```



```
|           ENABLE SEMANTIC SEARCH (optional)
|_
| # After docs are generated, setup RAG for AI-powered search:
| /docs:rag                         ← One command does it all!
| |
|     |
|     └─ Chunks your docs into searchable pieces
|     └─ Creates database migration (pgvector)
|     └─ Generates OpenAI embeddings
|     └─ Outputs sync.sql ready to import
```

Quick Commands

```
# Simple: Generate docs for one page
/docs:generate https://myapp.com/dashboard

# Flow: Document a feature (user-focused by default)
/docs:flow "how to create employees"
/docs:flow "import payments" --technical      # Developer docs

# Batch: Document entire project
/docs:discover && /docs:plan && /docs:execute

# Maintain: Update after code changes
/docs:update --base main

# RAG: Setup semantic search for your docs
/docs:rag
```

CLI Commands

`aidocs init [PROJECT_NAME]`

Initialize the docs module in a project.

```
aidocs init .                      # Current directory
aidocs init my-project              # New directory
aidocs init . --force                # Overwrite existing
aidocs init . --ai cursor            # Use with Cursor
```

Options:

Option	Description
`--ai`	AI assistant: `claude`, `cursor`, `copilot` (default: `claude`)
`--force, -f`	Overwrite existing files
`--no-git`	Skip git initialization

`aidocs check`

Check for required tools and dependencies.

```
aidocs check
```

`aidocs version`

Show version information.

`aidocs update`

Update aidocs to the latest version.

```
aidocs update          # Update from PyPI
aidocs update --github # Update from GitHub (latest)
```

Options:

Option	Description
--github	Install latest from GitHub instead of PyPI

Automatically detects and uses the appropriate package manager (uv, pipx, or pip).

`aidocs rag-chunks`

Chunk markdown files for vector database import.

```
aidocs rag-chunks          # Chunk all files in docs/
aidocs rag-chunks docs/users # Chunk specific directory
aidocs rag-chunks --force    # Re-chunk all files
aidocs rag-chunks --dry      # Preview only
```

Options:

Option	Description
`--force, -f`	Re-chunk all files (ignore cache)
`--dry`	Preview without writing files

What it does:

1. Scans directory for `.md` files
2. Splits at `##` headings into chunks
3. Creates `.chunks.json` files alongside each `.md`
4. Maintains `docs/.chunks/manifest.json` for change tracking

Output structure:

```
docs/
└── users/
    ├── lifecycle.md
    └── lifecycle.chunks.json      # Chunks for this file
└── campaigns/
    ├── lifecycle.md
    └── lifecycle.chunks.json
└── .chunks/
    └── manifest.json            # Tracking file
```

Next step: Run `aidocs rag-vectors` to generate embeddings

`aidocs rag-vectors`

Generate embeddings and SQL for vector database import.

```
aidocs rag-vectors          # Generate embeddings and SQL
aidocs rag-vectors --dry     # Preview what would be synced
```

```
aidocs rag-vectors --force          # Re-sync all files
aidocs rag-vectors --table my_docs # Custom table name
```

Options:

Option	Description
`--force, -f`	Re-sync all files (ignore last sync)
`--dry`	Preview without generating embeddings
`--table, -t`	Target table name (default: `doc_embeddings`)

Requires: `OPENAI_API_KEY` environment variable**What it does:**

1. Reads chunk files from `docs/.chunks/`
2. Calls OpenAI API to generate embeddings (text-embedding-3-small)
3. Creates `docs/.chunks/sync.sql` with INSERT statements
4. Tracks sync state to avoid re-processing unchanged files

Output: `docs/.chunks/sync.sql`

```
BEGIN;
INSERT INTO doc_embeddings (file_path, content, chunk_index, title, metadata,
embedding)
VALUES ('docs/users/lifecycle.md', '...', 0, 'Overview', '{...}'::jsonb, '[0.001,
...]'::vector);
-- ... more inserts
COMMIT;
```

Import to database:

```
psql $DATABASE_URL -f docs/.chunks/sync.sql
```

Slash Commands

After running `aidocs init`, these commands are available in Claude Code:

Command	Description	Requires Playwright
`/docs:init`	Configure project settings, credentials, output style	No
`/docs:generate`	Generate docs for a single page with screenshots	Yes
`/docs:analyze`	Analyze codebase for a route (no browser)	No
`/docs:batch`	Generate docs for multiple pages	Yes
`/docs:update`	Update docs based on git diff	Optional
`/docs:discover`	Scan codebase, discover all modules	No
`/docs:plan`	Create ordered documentation plan	No
`/docs:execute`	Execute plan, generate all docs	Yes
`/docs:explore`	Interactive UI exploration with Playwright	Yes
`/docs:flow ""`	Document a feature with screenshots (use `--technical` for dev docs)	Optional
`/docs:rag-vectors`	Generate embeddings and SQL for vector DB import	No
`/docs:rag-init`	Generate database migration for vector embeddings	No
`/docs:rag`	Setup RAG: chunks → migration → embeddings (all-in-one)	No
`/docs:export-pdf`	Export markdown documentation to PDF with TOC	Yes (Playwright)

/docs:init`

Interactive setup wizard that:

- Detects your tech stack (Laravel, Vue, React, Next.js, etc.)
- Asks for project name, audience, and documentation tone
- Configures authentication method (file, env vars, or manual)
- Sets output directory and screenshot preferences

'/docs:generate'

Generate documentation for a single page:

```
/docs:generate https://myapp.com/campaigns          # Uses base URL from config  
/docs:generate /campaigns                          # With authentication  
/docs:generate /settings --auth user:pass
```

Features:

- Captures full-page screenshots
- Analyzes UI elements visually
- Searches codebase for related code
- Detects forms, buttons, and interactive elements
- Offers to document user flows step-by-step

'/docs:update'

Update existing documentation based on code changes:

```
/docs:update                      # Compare against main  
/docs:update --base staging        # Compare against staging branch  
/docs:update --dry-run             # Preview changes without applying  
/docs:update --screenshots        # Also refresh screenshots
```

What it does:

1. Gets git diff between current branch and base
2. Analyzes changed frontend/backend files
3. Maps code changes to affected features
4. Finds and updates related documentation
5. Optionally refreshes screenshots
6. Offers to stage/commit doc changes

Perfect for: Running before creating a PR to ensure docs stay in sync with code.

'/docs:analyze'

Analyze codebase without browser automation:

```
/docs:analyze /campaigns
/docs:analyze /api/users
```

`/docs:batch`

Generate documentation for multiple pages:

```
/docs:batch urls.txt          # From file
/docs:batch --discover --base-url https://myapp.com # Auto-discover routes
```

`/docs:discover`

Scan your codebase to discover all modules and their structure:

```
/docs:discover                  # Discover all modules
/docs:discover --dry            # Preview without saving
/docs:discover campaigns        # Analyze only one module
```

What it analyzes:

- Backend: Models, controllers, routes, validation rules
- Frontend: Pages, components, forms, state management
- Relationships: Foreign keys, ORM relationships, cross-module navigation

Creates `docs/.knowledge/` with:

```
docs/.knowledge/
├── _meta/
│   ├── project.json           # Project-level info
│   └── modules-index.json     # List of discovered modules
└── modules/
    ├── campaigns/
    │   ├── entity.json          # Fields, types, relationships
    │   ├── routes.json           # API endpoints
    │   ├── components.json       # UI components
    │   └── validation.json        # Validation rules
    └── users/
        └── ...
└── relationships/             # Cross-module relationships
```

Next step: Run `/docs:plan` to create documentation plan

'/docs:plan'

Create an ordered documentation plan based on discovered modules:

```
/docs:plan                      # Create plan interactively  
/docs:plan --auto                # Auto-generate plan (no prompts)  
/docs:plan --show                # Show existing plan
```

What it does:

1. Reads discovered modules from `docs/.knowledge/`
2. Analyzes dependencies and relationships
3. Suggests documentation order (core modules first)
4. Creates `docs/plan.yml` with the plan

Output: `docs/plan.yml`

```
modules:  
  - name: users  
    priority: 1  
    reason: "Core module - other modules depend on it"  
    document:  
      lifecycle: true  
      include_errors: true  
    status: pending  
  
  - name: campaigns  
    priority: 2  
    document:  
      lifecycle: true  
      flows:  
        - "duplicate campaign"  
    status: pending  
  
cross_module_flows:  
  - name: "user registration to first campaign"  
    modules: [users, campaigns]  
    status: pending
```

Next step: Run `/docs:execute` to generate documentation

'/docs:execute'

Execute the documentation plan and generate all docs:

```
/docs:execute                      # Execute full plan
/docs:execute --module campaigns    # Execute only one module
/docs:execute --continue            # Continue from where it stopped
/docs:execute --dry                 # Preview what would be generated
```

What it does:

1. Reads `docs/plan.yml`
2. For each module in order:
 - Runs `explore` (if needed)
 - Generates lifecycle documentation
 - Captures screenshots
 - Writes to `docs/{module}/`
3. Updates plan status as it progresses
4. Generates cross-module flows last

Output structure:

```
docs/
├── index.md                      # Auto-generated with links
└── users/
    ├── index.md                  # Module overview
    ├── lifecycle.md              # CRUD documentation
    ├── user-registration-to-campaign.md # Cross-module flow (first module)
    └── images/
└── campaigns/
    ├── index.md
    ├── lifecycle.md
    ├── duplicate-campaign.md    # Custom flow
    └── images/
```

Resume support: If execution stops, run `/docs:execute --continue` to resume

'/docs:explore'

Interactively explore a module's UI with Playwright:

```
/docs:explore campaigns          # Explore all campaign pages
/docs:explore users --page /users/create # Specific page
/docs:explore orders --depth deep    # Thorough exploration
```

What it discovers:

- Conditional fields (checkbox reveals more inputs)
- Validation messages (tries invalid data)
- UI state changes (what happens when you click)
- Cross-page effects (create here → appears there)

'/docs:flow'"'

Document a feature with screenshots and step-by-step instructions. By default, creates **user-focused** documentation. Use `--technical` for developer documentation.

```
/docs:flow "how to create employees"          # User guide with screenshots
/docs:flow "import payments from csv"        # User guide with screenshots
/docs:flow "payment processing" --technical   # Developer docs with code
/docs:flow "stripe webhooks" --technical       # Developer docs with code
/docs:flow "user registration" --no-screenshots # Skip screenshots
```

Arguments:

- `--technical` - Generate developer-focused documentation with code snippets
- `--no-screenshots` - Skip UI screenshot capture

Output modes:

Mode	Audience	Output
Default	End users	Screenshots, plain English, step-by-step guide
`--technical`	Developers	Code snippets, file paths, mermaid diagrams

Output: `docs/flows/{kebab-case-title}.md`

Example: User-focused (default)

How to Import Payments

Import payment records from a CSV file.

Before You Start

- Prepare a CSV with columns: date, amount, description
- Maximum 10,000 rows per import

Steps

Step 1: Go to Payroll

Navigate to **Payroll** from the sidebar.

![Payroll Page](./images/payroll-page.png)

Step 2: Click Import

Click the **Import Payments** button.

![Import Button](./images/import-button.png)

Step 3: Upload Your File

Select your CSV file and click **Start Import**.

What Happens Next

- Import runs in background
- You'll receive an email when complete

Example: Technical (--technical)

Import Payments Flow

Architecture

sequenceDiagram: User → Controller → Job → Database

Entry Points

```
| Trigger | Route |
|-----|-----|
| UI | POST /payroll/import |
| CLI | php artisan payments:import |
```

Execution Flow

File: `app/Http/Controllers/PayrollController.php:45`
public function import(Request \$request) { ... }

File: `app/Jobs/ImportPaymentsJob.php:28`
public function handle() { ... }

Screenshots require:

- Playwright MCP installed
- `urls.base` configured in `docs/config.yml`

'/docs:rag-vectors'

Generate embeddings and SQL for syncing documentation to a PostgreSQL vector database.

```
/docs:rag-vectors          # Generate sync SQL (smart)
/docs:rag-vectors --dry    # Preview what would be synced
/docs:rag-vectors --force  # Re-sync all files
```

Prerequisites:

- Run `aidocs rag-chunks` first to create chunk files
- Set `OPENAI_API_KEY` environment variable

What it does:

1. Reads chunk files from `docs/.chunks/manifest.json`
2. Compares against last sync to find changes
3. Generates embeddings via OpenAI API (only for new/changed chunks)
4. Creates `docs/.chunks/sync.sql` with INSERT statements

Smart sync:

- Unchanged files → Skip (no API calls)
- Changed files → Re-generate embeddings
- New files → Generate embeddings
- Deleted files → Add DELETE statements

Output:

```
📊 Sync Summary:
Unchanged: 12 files (skipped)
Changed: 2 files (8 chunks)
New: 1 file (3 chunks)

📄 Generated: docs/.chunks/sync.sql
```

Run with:

```
psql $DATABASE_URL -f docs/.chunks/sync.sql
```

`/docs:rag-init`

Generate a database migration for storing documentation embeddings with pgvector.

```
/docs:rag-init          # Default: 1536 dimensions
/docs:rag-init --dimensions 3072 # For text-embedding-3-large
/docs:rag-init --table my_docs   # Custom table name
```

What it does:

1. Detects your framework (Laravel, Prisma, TypeORM, Drizzle, Django)
2. Generates the appropriate migration file
3. Creates table with pgvector support for similarity search

Supported Frameworks:

Framework	Detection	Output
Laravel	`composer.json`	PHP migration with `'\$table->vector()'`
Prisma	`schema.prisma`	Prisma schema addition
TypeORM	`package.json`	TypeScript migration class
Drizzle	`drizzle-orm`	Schema + SQL migration
Django	`manage.py`	Django migration with pgvector
Fallback	None detected	Raw PostgreSQL SQL

Table Structure:

```
doc_embeddings
└── id           UUID PRIMARY KEY
└── file_path    VARCHAR(500)      # Path to .md file
└── content      TEXT             # Document content
└── chunk_index  INTEGER         # For large docs split into chunks
└── title        VARCHAR(255)     # Document title
```

```

└── metadata      JSONB          # Tags, module, category, etc.
└── embedding    VECTOR(1536)   # OpenAI embedding
└── created_at   TIMESTAMP
└── updated_at   TIMESTAMP

```

Indexes:

- `file_path` - B-tree index for path lookups
- `embedding` - HNSW index for fast vector similarity search

Requirements:

- PostgreSQL with [pgvector](<https://github.com/pgvector/pgvector>) extension

Example workflow:

```

# 1. Generate migration
/docs:rag-init

# 2. Run migration
php artisan migrate           # Laravel
npx prisma migrate dev        # Prisma
python manage.py migrate       # Django

# 3. Chunk your docs
aidocs rag-chunks

# 4. Generate embeddings and sync
aidocs rag-vectors

```

`/docs:rag`

The easy way - Setup RAG (Retrieval Augmented Generation) for your documentation in one command:

```

/docs:rag                  # Full setup
/docs:rag --skip-migration # Skip migration (table already exists)
/docs:rag --force          # Re-chunk and re-sync everything
/docs:rag --dry             # Preview what would happen

```

What it does automatically:

1. Checks/creates documentation chunks (`aidocs rag-chunks`)
2. Generates database migration (`/docs:rag-init`)

3. Prompts you to run the migration
4. Generates embeddings and SQL (`aidocs rag-vectors`)

Output:

```

✓ RAG Setup Complete!

📊 Summary:
Documentation files: 8
Chunks created: 24
Embeddings generated: 24

📄 Files created:
✓ docs/.chunks/manifest.json
✓ database/migrations/..._create_doc_embeddings_table.php
✓ docs/.chunks/sync.sql

🚀 Final step:
psql $DATABASE_URL -f docs/.chunks/sync.sql

```

Requirements:

- PostgreSQL with [pgvector](https://github.com/pgvector/pgvector) extension
- `OPENAI_API_KEY` environment variable

`'/docs:export-pdf'`

Export markdown documentation to PDF with auto-generated table of contents using Playwright MCP.

```
/docs:export-pdf docs/pages/dashboard.md          # Export single file
/docs:export-pdf docs/flows-sync-users.md --output manual.pdf # Custom filename
```

What it does:

1. Reads the markdown file
2. Extracts H1/H2 headings to build a clickable table of contents
3. Converts markdown to styled HTML (code blocks, tables, images)
4. Uses Playwright MCP to render and export as PDF
5. Saves to `docs/exports/` directory

Output: `docs/exports/{filename}.pdf`

Features:

- Auto-generated TOC from H1/H2 headings with clickable links
- PDF-friendly styling (page breaks at H1, code block formatting)
- Embedded images (converted to base64)
- A4 format with proper margins

Example:

The screenshot shows a software interface for exporting a Markdown file. At the top, there's a progress bar indicating the export status. Below it, a tree view shows the contents of the file: a Table of Contents with sections like 'Dashboard Overview', 'Key Metrics', 'Navigation', 'Components', and 'Configuration'. Under 'Rendering PDF...', it specifies 'Format: A4' and 'Pages: 5'. At the bottom, a green checkmark indicates 'PDF exported!' and provides a link to the generated file: 'docs/exports/dashboard.pdf (245 KB)'.

Requirements:

- Playwright MCP must be available

Knowledge Base

The intelligent commands build a `docs/.knowledge/` folder:

```
docs/.knowledge/
├── _meta/                      # Project info
└── modules/
    ├── campaigns/
    │   ├── entity.json        # Entity definition
    │   ├── routes.json         # API routes
    │   ├── validation.json     # Validation rules
    │   ├── flows/              # User flows
    │   └── ui-states/          # Conditional UI
    └── users/
        └── ...
```

```

└── relationships/          # Cross-module relationships
└── cross-module-flows/    # Flows spanning modules

```

This knowledge powers smarter documentation generation.

Intelligent Workflow

For Single Flow (Quick)

```

/docs:flow "sync users from discord"      → Analyzes code, generates docs with diagrams
/docs:flow "import payments from csv"     → Includes UI screenshots if route detected

```

For Entire Project (Batch)

```

/docs:discover           → Scans codebase, finds all modules
↓
/docs:plan               → Creates ordered documentation plan
↓
/docs:execute            → Generates all docs with screenshots

```

Example Session

```

# Option A: Document a specific flow
/docs:flow "sync users from discord"
/docs:flow "import payments from csv"
/docs:flow "how stripe webhooks work"

# Option B: Document entire project
/docs:discover
/docs:plan
/docs:execute

# Resume if interrupted
/docs:execute --continue

# After code changes
/docs:update --base main

# Backend integration
# Import with UI screenshots
# Webhook handling

# Find all modules
# Create plan (docs/plan.yml)
# Generate all documentation

```

What Makes It Smart

Capability	How It Works
Conditional UI	Clicks checkboxes/toggles, observes what fields appear
Validation Discovery	Submits empty/invalid forms, captures error messages
Cross-Page Tracking	Creates data, verifies it appears in lists/dashboards
Entity Lifecycle	Documents full create → view → edit → delete flow
Modular Analysis	One module at a time, scales to large projects
Code + UI Correlation	Matches frontend components to backend validation

Configuration

After running `/docs:init`, a `docs/config.yml` is created:

```
project:
  name: "My App"
  type: saas

style:
  tone: friendly # friendly | professional | technical | minimal

urls:
  base: "https://myapp.com"

auth:
  method: file # file | env | manual

output:
  directory: ./docs
```

Authentication Methods

Method	Description
`file`	Credentials stored in `docs/auth` (gitignored)
`env`	Read from `DOCS_AUTH_USER` and `DOCS_AUTH_PASS`

`manual`	Pass `--auth user:pass` each time
----------	-----------------------------------

Output

Generated documentation includes:

- **Overview** - What the page is for
- **Features** - What users can do
- **Key Actions** - Buttons and actions explained
- **Screenshots** - Full-page captures
- **How-to Guides** - Step-by-step flows (optional)
- **Related Pages** - Navigation links

Requirements

- Python 3.11+
- Claude Code (or Cursor/Copilot)
- Playwright MCP (for browser-based commands)

Installing Playwright MCP

Add to your `~/.claude.json` or project `.mcp.json` :

```
{  
  "mcpServers": {  
    "playwright": {  
      "command": "npx",  
      "args": ["@anthropic/mcp-playwright"]  
    }  
  }  
}
```

Development

```
git clone https://github.com/binarcode/aidocs-cli.git  
cd aidocs-cli
```

```
uv venv && uv pip install -e .
aidocs check
```

License

MIT

Generated by aidocs export-pdf